

MACHINE LEARNING IN BIOINFORMATICS

ARTIFICIAL NEURAL NETWORKS

Philipp Benner

philipp.benner@bam.de


S.3 - eScience

Federal Institute for Materials Research and Testing (BAM)


July 22, 2023

- We will discuss a particular class of *Artificial Neural Networks (ANNs)* called *Multilayer Perceptrons (MLPs)*
- MLPs are feed-forward networks, i.e. without any loops (directed acyclic graphs)
- We motivate MLPs from logistic regression

DRAWINGS BY RAMON Y CAJAL (~ 1900)



A BIOLOGICAL NEURON




HODGKIN-HUXLEY MODEL

[HODGKIN AND HUXLEY, 1952]

$$I = C_m \frac{dV_m}{dt} + \bar{g}_K n^4 (V_m - V_K) + \bar{g}_{Na} m^3 h (V_m - V_{Na}) + \bar{g}_l (V_m - V_l)$$



$$\frac{dn}{dt} = \alpha_n(V_m)(1 - n) - \beta_n(V_m)n, \quad \frac{dm}{dt} = \alpha_m(V_m)(1 - m) - \beta_m(V_m)m$$

$$\frac{dh}{dt} = \alpha_h(V_m)(1 - h) - \beta_h(V_m)h$$



FIRST MATHEMATICAL MODELS

[MCCULLOCH AND PITTS, 1943]




FIRST MATHEMATICAL MODELS


[McCulloch AND PITTS, 1943]

- ▶ combining several McCulloch/Pitts neurons, we can build more complex boolean functions, e.g.:

- $y = \neg x_1 \vee (x_2 \wedge x_3)$:



- $y = XOR(x_1, x_2) = (x_1 \wedge \neg x_2) \vee (\neg x_1 \wedge x_2)$:




- ▶ Theorem: Networks of McCulloch/Pitts neurons can realize all boolean functions

- McCulloch-Pitts networks must be designed, i.e. there exists no training algorithm

Short and incomplete history of neural network models:

- McCulloch-Pitts networks [McCulloch and Pitts, 1943]
- The perceptron [Rosenblatt, 1957]
- Multilayer perceptrons [Werbos, 1974, McClelland et al., 1986]
- Backpropagation algorithm [LeCun et al., 1988]

ALEXNET - THE BREAKTHROUGH



- AlexNet [Krizhevsky et al., 2012] is the first neural network that performs better than traditional feature extraction and classification methods on images

FROM LOGISTIC REGRESSION TO ANNs

LOGISTIC REGRESSION


- We convert $x^\top \theta$ to probabilities

$$\text{pr}(Y = 1 | \theta) = \sigma(x^\top \theta)$$

- The function σ denotes the sigmoid function


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Visual representation of logistic regression:




where $x^{(j)}$ is the j -th value of the input vector $x \in \mathbb{R}^p$

GENERALIZING LOGISTIC REGRESSION




GENERALIZING LOGISTIC REGRESSION




- Instead of θ we use w for weights. In addition, q refers to the number of inputs or features, which for neural networks is much lower than the number of parameters p

GENERALIZING LOGISTIC REGRESSION




- Instead of θ we use w for weights. In addition, q refers to the number of inputs or features, which for neural networks is much lower than the number of parameters p
- We introduce multiple output neurons $\tilde{y}_1, \dots, \tilde{y}_k$

GENERALIZING LOGISTIC REGRESSION




- Instead of θ we use w for weights. In addition, q refers to the number of inputs or features, which for neural networks is much lower than the number of parameters p
- We introduce multiple output neurons $\tilde{y}_1, \dots, \tilde{y}_k$
- Weights are compiled into a single weight matrix $W \in \mathbb{R}^{p \times k}$

GENERALIZING LOGISTIC REGRESSION



- Instead of θ we use w for weights. In addition, q refers to the number of inputs or features, which for neural networks is much lower than the number of parameters p
- We introduce multiple output neurons $\tilde{y}_1, \dots, \tilde{y}_k$
- Weights are compiled into a single weight matrix $W \in \mathbb{R}^{p \times k}$
- We stack several layers, where each layer ℓ has its own weight matrix $W^{(\ell)}$ and activation function σ_ℓ

ARTIFICIAL NEURAL NETWORK



Feedforward Artificial Neural Network

A feedforward neural network consists of a variable number of layers. Each layer ℓ has its own weight matrix $W^{(\ell)} \in \mathbb{R}^{k_\ell \times k_{\ell+1}}$, where k_ℓ denotes the number of nodes within layer ℓ . The input layer has weight matrix $W^{(1)} \in \mathbb{R}^{k_1 \times k_2}$, where $q = k_1$ is the dimension of the input data. The activation function σ_ℓ is a layer-specific function and applied independently to each node.

ARTIFICIAL NEURAL NETWORK - NOTATION


- Let $X = X^{(1)} \in \mathbb{R}^{n \times q}$ denote a data matrix with n samples or observations, each of dimension q
- The output of the ℓ -th layer of a neural network is denoted

$$X^{(\ell+1)} = \sigma_\ell(X^{(\ell)}W^{(\ell)})$$


- The activation function σ_ℓ is applied to each cell of the matrix $X^{(\ell)}W^{(\ell)}$
- The prediction \tilde{y} of a neural network with L layers is the output $X^{(L+1)}$ of the last layer
- *Deep neural networks* have a large number of layers L

ACTIVATION FUNCTIONS

- Linear functions $\sigma(x) = ax + b$ are not very useful, except for output neurons¹
- Traditionally there were mainly two activation functions:



- Rectifier linear units (ReLU) [Glorot et al., 2011] and Softplus show better performance for training deep neural networks



¹Stacking of linear functions results in a linear function

ACTIVATION FUNCTIONS

- Especially important is the activation function of the output layer

- Classification:

- ▶ For binary classification problems we use one output node with sigmoid activation
- ▶ The softmax activation $\sigma : \mathbb{R}^k \rightarrow (0, 1)^k$ is used for multiclass problems

$$\sigma(x)_i = \frac{e^{x_i}}{\sum_{j=1}^k e^{x_j}}$$

- ▶ The output of the softmax is interpreted as class probabilities


- Regression:

- ▶ In most cases simply the identity function $\sigma(x) = x$
- ▶ ReLU can be used for only positive targets

ARCHITECTURES

- How do we choose the architecture? How many hidden layers? How many neurons in each layer?
- Neural networks with one hidden layer and an arbitrary amount of neurons are universal approximators [Hornik et al., 1989]
- Deep neural networks seem to perform better in practice than shallow networks with many neurons
- Each application domain has its own Architectures
 - ▶ Image processing: convolutional neural networks (CNNs)
 - ▶ Materials: Graph neural networks
 - ▶ Machine translation: transformers / attention

ANN INTERPRETATION



TRAINING ANNs

TRAINING ANNS

- Similar to logistic regression, the weights of ANNs are computed using gradient descent
- Gradient descent is an iterative algorithm to minimize a given loss function \mathcal{L}_W , i.e.

$$\hat{W} = \arg \min_W \mathcal{L}_W(X, y)$$

where (X, y) denotes the training data and $W = (W^{(1)}, W^{(2)}, \dots, W^{(L)})$ the weights of the neural network

- Logistic regression uses the negative log-likelihood as loss function
- What loss functions do we use for neural networks?

TRAINING ANNs - LOSS FUNCTIONS

- Let f_W denote the neural network with weights W and (X, y) a training data set

- Regression:

- ▶ Mean squared error (similar to OLS):

$$\mathcal{L}_W = \frac{1}{n} \|y - f_W(X)\|_2^2$$

- ▶ Mean absolute error:

$$\mathcal{L}_W = \frac{1}{n} \|y - f_W(X)\|_1$$

Useful when targets y contain outliers

TRAINING ANNs - LOSS FUNCTIONS

■ Classification:

- Binary cross-entropy ($y_i \in \{0, 1\}$):

$$\mathcal{L}_W = -\frac{1}{n} \sum_{i=1}^n y_i \log(f_W(x_i)) + (1 - y_i) \log(1 - f_W(x_i))$$

Equivalent to maximum likelihood of logistic regression

- Multiclass cross-entropy with k classes ($y_i \in \{0, 1\}^k$):

$$\mathcal{L}_W = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k y_{i,j} \log(f_W(x_i)_j) = -\frac{1}{n} \sum_{i=1}^n y_i^\top \log(f_W(x_i))$$

where $y_{i,j}$ is 1 if the i -th sample belongs to class j and $f_W(x_i)_j$ denotes the output of the j -th node of the last layer of the neural network f_W

■ The loss function is typically not convex!


TRAINING ANNS - GRADIENT DESCENT

Gradient descent

Gradient descent updates the weights at each iteration t according to the following update rule:

$$W_{t+1} = W_t - \gamma \nabla_{W_t} \mathcal{L}_{W_t}(X, y)$$

where γ is a parameter that determines the step size



- The larger γ the more likely the algorithm jumps over local optima, but the higher the chance of divergence

Stochastic gradient descent (SGD)

Stochastic gradient descent (SGD) selects at each iteration t a single training sample (x_t, y_t) with $t \in \{1, \dots, n\}$ at random and updates the weights based on this one sample:

$$W_{t+1} = W_t - \gamma \nabla_{W_t} \mathcal{L}_{W_t}(x_t, y_t)$$

- The stochastic nature of SGD can help to bypass local optima [Masters and Luschi, 2018]
- SGD is slow, because we have to update weights for each sample and cannot utilize parallel computation

Mini-batch SGD

Mini-batch SGD selects at each iteration t a random subset (X_t, y_t) of m samples $X_t = (x_{t_1}, \dots, x_{t_m})$, $y_t = (y_{t_1}, \dots, y_{t_m})$ with $t_k \in \{1, \dots, n\}$ and updates the weights accordingly:

$$W_{t+1} = W_t - \gamma \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)$$

(X_t, y_t) is called a mini-batch and m controls the mini-batch size

- Practice has shown that (mini-batch) SGD seems to improve generalization [Hoffer et al., 2017]
- We typically select $m = 32$ or $m = 64$

Gradient descent with momentum

The update rule of gradient descent with momentum is


$$M_{t+1} = \beta M_t + \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)$$

$$W_{t+1} = W_t - \gamma M_{t+1}$$

where γ is the usual step size and β determines the weight of previous updates

- For $\beta = 0$ we obtain vanilla gradient descent
- Intuitive explanation: The mass of a stone rolling downhill adds momentum
- In some cases gradient descent with momentum can help to jump out of local optima
- It makes SGD or mini-batch SGD more stable

TRAINING ANNs - GRADIENT DESCENT



- Gradient descent converges slowly when being stuck on plateaus
- The gradient itself is not very informative
- Some gradient descent methods only use the sign of the gradient (i.e. Rprop [Riedmiller and Braun, 1992])
- Rprop does not work for mini-batches

Root Mean Square Propagation (RMSProp)

RMSProp uses the following update rule:

$$\begin{aligned}V_{t+1} &= \beta V_t + (1 - \beta) (\nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t))^2 \\W_{t+1} &= W_t - \frac{\gamma}{\sqrt{V_{t+1}}} \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)\end{aligned}$$

where γ denotes the step size and β decay rate for averaging over previous gradients

■ For $\beta = 0$ we obtain

$$W_{t+1} = W_t - \frac{\gamma}{|\nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)|} \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)$$

i.e. we only consider the sign of the gradient

Adaptive Moment Estimation (Adam)

Adam uses the following update rule:

$$\begin{aligned} M_{t+1} &= \beta_1 M_t + (1 - \beta_1) \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t) & | \text{ momentum} \\ V_{t+1} &= \beta_2 V_t + (1 - \beta_2) \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)^2 & | \text{ adaptive } \gamma \\ \hat{M}_{t+1} &= M_{t+1} / (1 - \beta_1^{t+1}), \quad \hat{V}_{t+1} = V_{t+1} / (1 - \beta_2^{t+1}) & | \text{ bias correction} \\ W_{t+1} &= W_t - \frac{\gamma}{\sqrt{\hat{V}_{t+1} + \epsilon}} \hat{M}_{t+1} \end{aligned}$$

where γ denotes the step size and β decay rate for averaging over previous gradients

- $M_0 = 0$ and $V_0 = 0$, hence early estimates are biased towards zero, \hat{M} and \hat{V} correct this bias

Adaptive Moment Estimation (Adam)

Adam uses the following update rule:

$$\begin{aligned}M_{t+1} &= \beta_1 M_t + (1 - \beta_1) \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t) && | \text{ momentum} \\V_{t+1} &= \beta_2 V_t + (1 - \beta_2) \nabla_{W_t} \mathcal{L}_{W_t}(X_t, y_t)^2 && | \text{ adaptive } \gamma \\\hat{M}_{t+1} &= M_{t+1} / (1 - \beta_1^{t+1}), \quad \hat{V}_{t+1} = V_{t+1} / (1 - \beta_2^{t+1}) && | \text{ bias correction} \\W_{t+1} &= W_t - \frac{\gamma}{\sqrt{\hat{V}_{t+1} + \epsilon}} \hat{M}_{t+1}\end{aligned}$$

where γ denotes the step size and β decay rate for averaging over previous gradients

- Adam is the default for training neural networks
- RMSProp with bias-correction and momentum

VANISHING AND EXPLODING GRADIENT PROBLEM

- Suppose we want the gradient of the first weight matrix of a neural network with L layers
- The derivative of the first weight matrix is defined by the *chain rule*:

$$\frac{\partial}{\partial W^{(1)}} \mathcal{L}_W(X, y) = \frac{\partial \mathcal{L}_W}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial X^{(L)}} \frac{\partial X^{(L)}}{\partial X^{(L-1)}} \cdots \frac{\partial X^{(2)}}{\partial W^{(1)}}(X, y)$$

- We use the partial derivative $\partial/\partial W^{(1)}$ to denote that all other weight matrices

$$W^{(2)}, W^{(3)}, \dots, W^{(L)}$$

are treated as constants at their current values


- In deep neural networks, the gradient of the first layers will behave very differently than the gradient of the last layers

VANISHING AND EXPLODING GRADIENT PROBLEM

- In deep neural networks, the gradient of the first layers will behave very differently than the gradient of the last layers
- For the first layers, the gradient easily vanishes or explodes
- Strategies to counter this problem
 - ▶ Normalization of training data
 - ▶ Proper weight initialization
 - ▶ Activation functions such as ReLU
 - ▶ Normalization of layer outputs $X^{(\ell)}$ (batch normalization)
 - ▶ Skip connections: $X^{(\ell+1)} = \sigma_\ell(X^{(\ell)}W^{(\ell)}) + X^{(\ell)}$

VANISHING AND EXPLODING GRADIENT PROBLEM

- Sigmoid and Tanh activations cause vanishing gradients:




- ReLU is less prone to vanishing gradients
- ReLU often leads to inactive nodes during training (zero gradient)
- LeakyReLU is often used instead

TRAINING ANNs - REMARKS

- How do we compute the gradient for training neural networks?
- Backpropagation algorithm for neural networks
- Libraries such as Tensorflow or PyTorch use automatic differentiation (AD)
- With AD it is possible to compute the gradient of arbitrary functions, including neural networks
- Developing deep neural networks requires much testing (watch the gradient!)
 - ▶ Weights must stay within a reasonable range
 - ▶ Training data must be normalized
 - ▶ Local optima must be bypassed using SGD
 - ▶ Overfitting must be prevented

TRAINING ANNs - REMARKS

- Early stopping prevents overfitting
- We use a small portion of the training data for validation






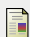
■ Statistics:

- ▶ Define the optimization problem (e.g. maximum likelihood)
- ▶ Use explicit regularization, especially for overparameterized models
- ▶ Check that there is a unique solution
- ▶ Develop numerical methods for finding the solution





■ Modern machine learning:

- ▶ Define the optimization problem (e.g. minimize MSE) using an overparameterized model
- ▶ Use a variety of tricks to compute solutions that generalize well
- ▶ The gradient method itself is considered a method for regularization





REFERENCES I

-  GLOROT, X., BORDES, A., AND BENGIO, Y. (2011).
DEEP SPARSE RECTIFIER NEURAL NETWORKS.
In Proceedings of the fourteenth international conference on artificial intelligence and statistics, pages 315–323. JMLR Workshop and Conference Proceedings.
-  HODGKIN, A. L. AND HUXLEY, A. F. (1952).
A QUANTITATIVE DESCRIPTION OF MEMBRANE CURRENT AND ITS APPLICATION TO CONDUCTION AND EXCITATION IN NERVE.
The Journal of physiology, 117(4):500.
-  HOFFER, E., HUBARA, I., AND SOUDRY, D. (2017).
TRAIN LONGER, GENERALIZE BETTER: CLOSING THE GENERALIZATION GAP IN LARGE BATCH TRAINING OF NEURAL NETWORKS.
Advances in neural information processing systems, 30.
-  HORNIK, K., STINCHCOMBE, M., AND WHITE, H. (1989).
MULTILAYER FEEDFORWARD NETWORKS ARE UNIVERSAL APPROXIMATORS.
Neural networks, 2(5):359–366.

REFERENCES II

-  KRIZHEVSKY, A., SUTSKEVER, I., AND HINTON, G. E. (2012).
IMAGENET CLASSIFICATION WITH DEEP CONVOLUTIONAL NEURAL NETWORKS.
Advances in neural information processing systems, 25.
-  LECUN, Y., TOURESKY, D., HINTON, G., AND SEJNOWSKI, T. (1988).
A THEORETICAL FRAMEWORK FOR BACK-PROPAGATION.
In *Proceedings of the 1988 connectionist models summer school*, volume 1, pages 21–28.
-  MASTERS, D. AND LUSCHI, C. (2018).
REVISITING SMALL BATCH TRAINING FOR DEEP NEURAL NETWORKS.
arXiv preprint arXiv:1804.07612.
-  MCCLELLAND, J. L., RUMELHART, D. E., AND HINTON, G. E. (1986).
THE APPEAL OF PARALLEL DISTRIBUTED PROCESSING.
MIT Press, Cambridge MA, pages 3–44.

REFERENCES III

-  MCCULLOCH, W. S. AND PITTS, W. (1943).
A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY.
The bulletin of mathematical biophysics, 5(4):115–133.
-  RIEDMILLER, M. AND BRAUN, H. (1992).
RPROP - A FAST ADAPTIVE LEARNING ALGORITHM.
In *Proc. of ISCIS VII*, Universitat. Citeseer.
-  ROSENBLATT, F. (1957).
THE PERCEPTRON, A PERCEIVING AND RECOGNIZING AUTOMATON PROJECT PARA.
Cornell Aeronautical Laboratory.
-  WERBOS, P. (1974).
BEYOND REGRESSION:" NEW TOOLS FOR PREDICTION AND ANALYSIS IN THE BEHAVIORAL SCIENCES.
Ph. D. dissertation, Harvard University.